

0.0.1 Amazon Athena

Amazon Athena[1] allows you to analyze data stored in Amazon S3 using SQL. It was created with the purpose of allowing anyone with SQL skills to quickly analyze large datasets. Athena will allow a user to run ad hoc SQL queries without the need to load or aggregate the data outside of S3 and can process structured, semi-structured and unstructured data sets. It is serverless so there is no need to deal with the setup or managing of infrastructure. This also allows Athena to scale automatically so the results continue to be fast with large datasets and complex queries. Athena uses Presto which is an open source SQL query engine which is optimized for low latency ad hoc querying. You can access Athena in multiple ways including the AWS Management Console, API or JDBC driver. It also integrates into Amazon Quicksight allowing you to visualize the data stored in your S3 environment based on your Athena queries. Athena is ideal for quick ad hoc querying, but can also handle complex joins, window functions and arrays. You only pay for the queries you run.

References

- [1] AWS. *Athena FAQs*. Web. URL: <https://aws.amazon.com/athena/faqs/> (cited on page 1).

Index

Amazon
Athena, [1](#)
Athena
Amazon, [1](#)